

Biology Refusal Classifier - Write-up

Data

Positives (“refuse”). 1,273 WMDP-bio questions from `cais/wmdp`

Negatives (“don’t refuse”). I used benign **MMLU biology and medical subsets** (`high_school_biology`, `college_biology`, `anatomy`, `virology`, `medical_genetics`, plus five clinical subjects). The idea is to make the model separate hazardous biology from ordinary biology. 2,094 negatives were down-sampled to 1,273 (equal fraction per subject); stratified 80/20 train-test split.

Models and results

Baseline: frozen embeddings + logistic regression. Each prompt is encoded once by `all-MiniLM-L6-v2` (a 6-layer, 22.7M-parameter sentence-transformer) into a 384-d vector. An L2-penalised logistic regression is fit on top: 385 trained parameters (384 weights + bias). This isolates how much signal is already linearly separable in general-purpose embedding space.

DistilBERT: full fine-tuning. `distilbert-base-uncased` with a new 2-class head (a 768x768 `pre_classifier` plus a 768x2 classifier, ~0.6M new parameters on top of the first-token representation). All ~67M parameters are updated: the 6 transformer layers, the embeddings, and the head; nothing is frozen. With only ~2k training examples this is cheap (trained in ~75 s on my old desktop computer, which just has a GTX 1080). Training: 3 epochs, AdamW at lr 2e-5 with linear decay, batch size 16, weight decay 0.01, max sequence length 256, fp32. `enable_full_determinism` fixes the run.

Headline metrics are shown right; AUPRC is for the refuse (positive) class.

Per-class precision / recall / F1 at the 0.5 threshold: DistilBERT: refuse

0.976 / 0.973 / 0.975, don’t-refuse 0.972 / 0.976 / 0.974; baseline: refuse

0.943 / 0.965 / 0.954, don’t-refuse 0.963 / 0.940 / 0.951 (raw counts in the

Figure 1 confusion matrices). Fine-tuning buys ~2.2 accuracy points. The

threshold-free metrics matter more for a safety filter: DistilBERT’s AUPRC of 0.995 means it holds near-perfect precision across almost the whole recall range, so the high-recall operating point a filter wants is nearly free.

Failure modes (DistilBERT)

Figure 3 shows sampled prompts from each confusion-matrix quadrant.

False negatives - missed refusals (7). Anecdotally, these WMDP positives look like ordinary biology.

False positives - benign flagged (6). Possibly benign MMLU questions tripping on hazard-adjacent vocabulary: “*Who discovered the first virus?*”, “*Nipah is a zoonotic paramyxovirus; where did it originate?*” If there are too many of these, the filter falls foul of Avid’s Rule of Usability - “security at the expense of usability comes at the expense of security” - because over-refusing benign biology drives users toward unfiltered alternatives.¹

Errors are often confident (Figure 4). DistilBERT is well-calibrated in aggregate (ECE 1.6%, Brier 0.021), but its mistakes split into two populations: ~half (6/13) sit near the 0.5 boundary, while the rest are confident - 5/13 at ≥ 0.99 , mean error-confidence 0.87. A low-confidence “escalate to review” band therefore catches 6/13 errors for only ~2% review load, then hits a wall: the confident errors need better data, not a threshold.

One systematic confounder (Figure 2): WMDP are uniformly short, while MMLU’s clinical-vignette questions give the benign class a long right tail, so prompt length weakly correlates with the label and could be learned as a shortcut. (I have not done any analysis to try to falsify this idea).

Out-of-domain behaviour

The classifier only ever sees biology in training, but a deployed filter mostly sees non-bio traffic. I scored it on 12 benign non-bio MMLU subjects (2,132 prompts, inference-only, not in training or the main metrics; Figure 5). Aggregate false-positive rate is a reassuring 2.3%, and pure humanities/STEM (history, philosophy, maths, physics, CS) are ~0%. But it is very uneven, with two categories where it does particularly badly: `us_foreign_policy` is flagged at 27% and `college_chemistry` at 9%. The former is really weird, and deserves some more investigation. WMDP = *Weapons of Mass Destruction*, so maybe the model learned something like “big things which could appear in the news”, which would include pathogenic viruses and foreign policy? But without more analysis this is pure speculation on my part.

What I’d do next

- **Better positives.** I don’t think the WMDP-bio question set represents a very realistic set of attacks. Many read pretty benign to me.
- **Threshold + calibration.** Pick a high-recall operating point from the PR curve rather than the 0.5 default; calibrate the (often over-confident) scores. Could add a low-confidence “escalate to review”/“escalate to a heavier model”.
- **Stricter evaluation.** Length-match the classes so prompt length can’t be used as a shortcut; test on a held-out *different* benign corpus. I’m worried the model could just learn “MMLU vs WMDP” style.
- **Long inputs.** Inputs are unrealistically short. Maybe windowing + max-pool aggregation, but I suspect this would lose a lot of nuance. Maybe a bigger input?

AI assistance

Claude Code under supervision helped quite a bit.

Figures




Figure 1. Confusion matrices on the held-out test set (n=503).




Figure 2. Prompt-length distribution by class. WMDP-bio (refuse) is uniformly short; MMLU (don't refuse) has a long right tail.



Figure 3. Sampled prompts from each confusion-matrix quadrant (DistilBERT), with the model's refuse-score. Green = correct, red = error.




Figure 4. Refuse-score distribution per confusion-matrix quadrant (DistilBERT). Histogram bars carry the large correct quadrants; the rug below each axis marks every individual prediction. Error quadrants (red) show many predictions far from the 0.5 boundary - confident mistakes.




Figure 5. Out-of-domain probe: refuse-score for every benign prompt across 12 non-bio MMLU subjects (inference-only). Red = false positive. Subjects ordered by false-positive rate.

1. “Avid’s Rule of Usability”, credited to Avi Douglan and widely cited on the Information Security Stack Exchange; I could not pin down the original post. [↩](#)